



Key Phrase Detection Using Compact CNN Accelerator IP

Reference Design

FPGA-RD-02066-1.0

October 2019

Disclaimers

Lattice makes no warranty, representation, or guarantee regarding the accuracy of information contained in this document or the suitability of its products for any particular purpose. All information herein is provided AS IS and with all faults, and all risk associated with such information is entirely with Buyer. Buyer shall not rely on any data and performance specifications or parameters provided herein. Products sold by Lattice have been subject to limited testing and it is the Buyer's responsibility to independently determine the suitability of any products and to test and verify the same. No Lattice products should be used in conjunction with mission- or safety-critical or any other application in which the failure of Lattice's product could create a situation where personal injury, death, severe property or environmental damage may occur. The information provided in this document is proprietary to Lattice Semiconductor, and Lattice reserves the right to make any changes to the information in this document or to any products at any time without notice.

Contents

Acronyms in This Document	7
1. Introduction	8
1.1. Design Process Overview	8
2. Setting Up the Basic Environment	9
2.1. Tools and Hardware Requirements	9
2.1.1. Lattice Tools	9
2.1.2. Hardware	9
2.2. Setting Up the Linux Environment for Machine Training	10
2.2.1. Installing the CUDA Toolkit	10
2.2.2. cuDNN Installation	11
2.2.3. Installing the Anaconda and Python3	11
2.2.4. Installing TensorFlow v1.12	12
2.2.5. Installing the Python Package	13
3. Preparing the Dataset	15
3.1. Selecting the Dataset	15
3.2. Downloading the Dataset	15
3.3. Data Augmentation	15
4. Training the Machine	16
4.1. Training Code Structure	16
4.2. Neural Network Architecture	17
4.2.1. Neural Network Architecture	17
4.2.2. Key Phrase Detection Network Output	19
4.2.3. Training Code Overview	19
4.2.4. Training from Scratch and/or Transfer Learning	24
5. Creating Frozen File	29
5.1. Generating the Frozen (.pb) File	29
6. Creating Binary File with SensAI	31
7. Hardware Implementation	35
7.1. Top Level Information	35
7.1.1. Block Diagram	35
7.1.2. Overall Operational Flow	35
7.1.3. Core Customization	36
8. RTL Bitstream Generation	37
9. Programming the Key Phrase Detection Demo	39
9.1. Functional Description	39
9.2. Programming Key Phrase Detection Demo on iCE40 SPI Flash	39
9.3. Run iCE40 Key Phrase Detection Demo on Hardware	45
Technical Support Assistance	46
Revision History	47

Figures

Figure 1.1. Lattice Machine Learning Design Flow	8
Figure 2.1. Lattice Himax HM01B0 Upduino Shield Board	9
Figure 2.2. CUDA Repo Download	10
Figure 2.3. CUDA Repo Installation	10
Figure 2.4. Fetch Keys	10
Figure 2.5. Updated Ubuntu Package Repositories	10
Figure 2.6. CUDA Installation Completed	11
Figure 2.7. cuDNN Library Installation	11
Figure 2.8. Anaconda Package Download	11
Figure 2.9. Anaconda Installation	12
Figure 2.10. License Terms Prompt	12
Figure 2.11. Installation Path Confirmation	12
Figure 2.12. Launch/Initialize Anaconda Environment Installation Complete	12
Figure 2.13. Anaconda Environment Activation	12
Figure 2.14. TensorFlow Installation	13
Figure 2.15. TensorFlow Installation Confirmation	13
Figure 2.16. TensorFlow Installation Complete	13
Figure 2.17. Easydict Installation	13
Figure 2.18. Anaconda Env – librosa Installation	14
Figure 3.1. Convert Sample Rate Script Execution	15
Figure 4.1. Training Code Directory Structure	16
Figure 4.2. Training Code Flow Diagram	19
Figure 4.3. Code Snippet – Model Settings	20
Figure 4.4. Code Snippet – Audio Processor	20
Figure 4.5. Code Snippet – Quantization Parameters Setting	22
Figure 4.6. Code Snippet – Forward Graph Fire Layers	23
Figure 4.7. Code Snippet – Convolution Quantization	23
Figure 4.8. Code Snippet – Training Loop	23
Figure 4.9. Code Snippet – Train Model	24
Figure 4.10. Code Snippet – Save Checkpoints	24
Figure 4.11. Code Snippet – Config File config.sh	25
Figure 4.12. Key Phrase Detection – Trigger Training with Default Options (Phase 1)	25
Figure 4.13. Key Phrase Detection – Training from Scratch (Phase 2)	26
Figure 4.14. Key Phrase Detection – Trigger Training with Transfer Learning	26
Figure 4.15. Key Phrase Detection – Training Logs	27
Figure 4.16. TensorBoard – Launch	27
Figure 4.17. TensorBoard – Link Default Output in Browser	27
Figure 4.18. Checkpoint Storage Directory Structure	28
Figure 5.1. Generating .pbtxt For Inference	29
Figure 5.2. Generated .pbtxt for Inference	29
Figure 5.3. Run genpb.py to Generate Inference .pb	29
Figure 5.4. Frozen Inference .pb Output	30
Figure 6.1. SensAI Home Screen	31
Figure 6.2. SensAI –Network File Selection	32
Figure 6.3. SensAI –Image Data File Selection	32
Figure 6.4. SensAI – Project Settings	33
Figure 6.5. SensAI – Analyze Project	33
Figure 7.1. Top Level Block Diagram Key Phrase Detection iCE40	35
Figure 8.1. Lattice Radiant Software	37
Figure 8.2. Lattice Radiant Software – Open Project	37
Figure 8.3. Lattice Radiant Software – Bitstream Generation	38
Figure 8.4. Lattice Radiant Software – Bitstream Generation Export Reports	38

Figure 9.1. iCE40 Key Phrase Demo Diagram	39
Figure 9.2. Radiant Programmer – Creating New Project	39
Figure 9.3. Radiant Programmer – Initial Project Window	40
Figure 9.4. Radiant Programmer – iCE40 UltraPlus Device Family Selection	40
Figure 9.5. Radiant Programmer – iCE40 UltraPlus Device Selection	41
Figure 9.6. Radiant Programmer – Bitstream Flashing Settings	42
Figure 9.7. Radiant Programmer – Filter-Firmware Bin File Flashing Setting	43
Figure 9.8. Radiant Programmer – Firmware Bin File Flashing Setting	44
Figure 9.9. Camera and LED Location	45

Tables

Table 4.1. Key Phrase Detection Training Network Topology17

Table 7.1. CNN Output Format35

Table 7.2. Core Parameters36

Acronyms in This Document

A list of acronyms used in this document.

Acronym	Definition
CKPT	Checkpoint
CNN	Convolutional Neural Network
EVDK	Embedded Vision Development Kit
FPGA	Field-Programmable Gate Array
LED	Light-emitting diode
MLE	Machine Learning Engine
NN	Neural Network
NNC	Neural Network Compiler
SD	Secure Digital
SDHC	Secure Digital High Capacity
SDXC	Secure Digital extended Capacity
SPI	Serial Peripheral Interface
VIP	Video Interface Platform
USB	Universal Serial Bus

1. Introduction

This document describes the key phrase detection design process using an iCE40 UltraPlus™ FPGA platform (HiMax HM01B0 UPduino Shield).

1.1. Design Process Overview

The design process involves the following steps:

1. Training the model
 - Setting up the basic environment
 - Preparing the dataset
 - Training the machine
 - Training the machine and creating the checkpoint data
 - Creating the frozen file (*.pb)
2. Compiling Neural Network
 - Creating the filter and firmware binary files with Lattice SensAI 2.1 program
3. FPGA design
 - Creating the FPGA Bitstream file
4. FPGA Bitstream and Quantized Weights and Instructions
 - Flashing the binary and bitstream files to iCE40 UPduino hardware

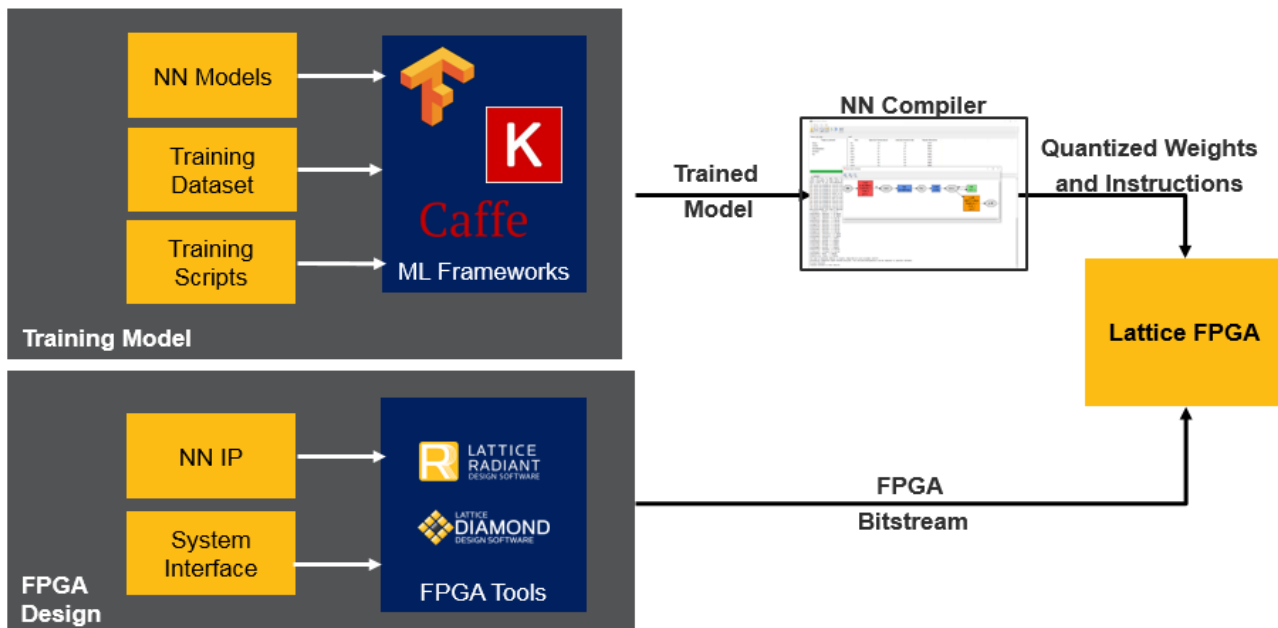


Figure 1.1. Lattice Machine Learning Design Flow

2. Setting Up the Basic Environment

2.1. Tools and Hardware Requirements

This section describes the required tools and environment setup for training and model freezing.

2.1.1. Lattice Tools

- Lattice Radiant Tool v1.1 – Refer to <http://www.latticesemi.com/latticeradiant>
- Lattice Radiant Programmer v1.1 – Refer to <http://www.latticesemi.com/latticeradiant>
- Lattice SensAI Compiler v2.1 – Refer to <https://www.latticesemi.com/Products/DesignSoftwareAndIP/AI/ML/NeuralNetworkCompiler>

2.1.2. Hardware

This design uses the HiMax HM01B0 Upduino Shield as shown in Figure 2.1. Refer to <http://www.latticesemi.com/en/Products/DevelopmentBoardsAndKits/HimaxHM01B0>.




Figure 2.1. Lattice Himax HM01B0 Upduino Shield Board

2.2. Setting Up the Linux Environment for Machine Training

2.2.1. Installing the CUDA Toolkit

To install the NVIDIA CUDA toolkit, run the commands below:

1. Download the NVIDIA CUDA toolkit.

```
$ curl -O
https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1604/x86_64/cu
da-repo-ubuntu1604_10.1.105-1_amd64.deb

(base) sib:~/kishan$ curl -O https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1604/x86_64/cuda-repo-ubuntu1604_10.1.105-1_amd64.deb
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100 2832 100 2832    0     0    514      0  0:00:05 0:00:05 --:--:-- 584
(base) sib:~/kishan$
```

Figure 2.2. CUDA Repo Download

2. Install the deb package.

```
$ sudo dpkg -i ./cuda-repo-ubuntu1604_10.1.105-1_amd64.deb

(base) sib:~/kishan$ sudo dpkg -i ./cuda-repo-ubuntu1604_10.1.105-1_amd64.deb
Selecting previously unselected package cuda-repo-ubuntu1604.
(Reading database ... 288236 files and directories currently installed.)
Preparing to unpack .../cuda-repo-ubuntu1604_10.1.105-1_amd64.deb ...
Unpacking cuda-repo-ubuntu1604 (10.1.105-1) ...
Setting up cuda-repo-ubuntu1604 (10.1.105-1) ...
(base) sib:~/kishan$
```

Figure 2.3. CUDA Repo Installation

3. Proceed with the installation.

```
$ sudo apt-key adv --fetch-keys
http://developer.download.nvidia.com/compute/cuda/repos/ubuntu1604/x86_64/7fa
2af80.pub

(base) sib:~/kishan$ sudo apt-key adv --fetch-keys http://developer.download.nvidia.com/compute/cuda/repos/ubuntu1604/x86_64/7fa2af80.pub
Executing: gpg --ignore-time-conflict --no-options --no-default-keyring --homedir /tmp/tmp.oqotmKcGn0 --no-auto-check-trustdb --trust-model
ng /etc/apt/trusted.gpg --keyring /etc/apt/trusted.gpg.d/diesch-testing.gpg --keyring /etc/apt/trusted.gpg.d/george-edison55-cmake-3_x.gpg -
--fetch-keys http://developer.download.nvidia.com/compute/cuda/repos/ubuntu1604/x86_64/7fa2af80.pub
gpg: key 7FA2AF80: "cudatools <cudatools@nvidia.com>" not changed
gpg: Total number processed: 1
gpg:      unchanged: 1
```

Figure 2.4. Fetch Keys

```
$ sudo apt-get update
```

```
(base) sib:~/kishan$ sudo apt-get update
Ign http://dl.google.com stable InRelease
Ign http://archive.ubuntu.com trusty InRelease
Ign http://extras.ubuntu.com trusty InRelease
Hit https://deb.nodesource.com trusty InRelease
Ign http://archive.canonical.com precise InRelease
Hit http://ppa.launchpad.net trusty InRelease
```

Figure 2.5. Updated Ubuntu Package Repositories

```
$ sudo apt-get install cuda-9-0
```

```
(base) sib:~/kishan$ sudo apt-get install cuda-9-0
Reading package lists... Done
Building dependency tree
Reading state information... Done
```

Figure 2.6. CUDA Installation Completed

2.2.2. cuDNN Installation

To install the cuDNN:

1. Create Nvidia developer account in <https://developer.nvidia.com>.
2. Download cuDNN library in https://developer.nvidia.com/compute/machine-learning/cudnn/secure/v7.1.4/prod/9.0_20180516/cudnn-9.0-linux-x64-v7.1.
3. Execute the commands below to install cuDNN:

```
$ tar xvf cudnn-9.0-linux-x64-v7.1.tgz
$ sudo cp cuda/include/cudnn.h /usr/local/cuda/include
$ sudo cp cuda/lib64/libcudnn* /usr/local/cuda/lib64
$ sudo chmod a+r /usr/local/cuda/include/cudnn.h /usr/local/cuda/lib64/libcudnn*
```

```
k$ tar xvf cudnn-9.0-linux-x64-v7.1.tgz
cuda/include/cudnn.h
cuda/NVIDIA_SLA_cuDNN_Support.txt
cuda/lib64/libcudnn.so
cuda/lib64/libcudnn.so.7
cuda/lib64/libcudnn.so.7.1.4
cuda/lib64/libcudnn_static.a
k$ sudo cp cuda/include/cudnn.h /usr/local/cuda/include
k$ sudo cp cuda/lib64/libcudnn* /usr/local/cuda/lib64
k$ sudo chmod a+r /usr/local/cuda/include/cudnn.h /usr/local/cuda/lib64/libcudnn*
k$ _
```

Figure 2.7. cuDNN Library Installation

2.2.3. Installing the Anaconda and Python3

To install the Anaconda and Python 3:

1. Go to <https://www.anaconda.com/distribution/#download>.
2. Download the Python 3 version of Anaconda for Linux.

```
sib:~/kishan$ wget https://repo.anaconda.com/archive/Anaconda3-2019.03-Linux-x86_64.sh
--2019-04-18 11:34:54-- https://repo.anaconda.com/archive/Anaconda3-2019.03-Linux-x86_64.sh
Resolving repo.anaconda.com (repo.anaconda.com)... 104.16.131.3, 104.16.130.3, 2606:4700::6810:8303, ...
Connecting to repo.anaconda.com (repo.anaconda.com)|104.16.131.3|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 685906562 (654M) [application/x-sh]
Saving to: 'Anaconda3-2019.03-Linux-x86_64.sh'

100%[=====
```

Figure 2.8. Anaconda Package Download

3. Run the command below to install the Anaconda environment.

```
$ sh Anaconda3-2019.03-Linux-x86_64.sh
```

Note: Anaconda3-<version>-Linux-x86_64.sh version may vary based on the release.

```
sib:~/kishan$ sh Anaconda3-2019.03-Linux-x86_64.sh

Welcome to Anaconda3 2019.03

In order to continue the installation process, please review the license
agreement.
Please, press ENTER to continue
>>> _
```

Figure 2.9. Anaconda Installation

4. Accept the license.

```
Do you accept the license terms? [yes/no]
[no] >>> yes_
```

Figure 2.10. License Terms Prompt

5. Confirm the installation path. Follow the instructions onscreen to change the default path.

```
Do you accept the license terms? [yes/no]
[no] >>> yes

Anaconda3 will now be installed into this location:
/home/sibridge/anaconda3

- Press ENTER to confirm the location
- Press CTRL-C to abort the installation
- Or specify a different location below

[/home/sibridge/anaconda3] >>> /home/sibridge/kishan/anaconda3_
```

Figure 2.11. Installation Path Confirmation

6. After installation, enter **No** as shown in [Figure 2.12](#).

```
installation finished.
Do you wish the installer to initialize Anaconda3
by running conda init? [yes/no]
[no] >>> no_
```

Figure 2.12. Launch/Initialize Anaconda Environment Installation Complete

2.2.4. Installing TensorFlow v1.12

To install the TensorFlow v1.12:

1. Activate the conda environment by running the command below:

```
$ source <conda directory>/bin/activate
```

```
sib:~/kishan$ source anaconda3/bin/activate
(base) sib:~/kishan$ _
```

Figure 2.13. Anaconda Environment Activation

2. Install the TensorFlow by running the command below:

```
$ conda install tensorflow-gpu==1.12.0
```

```
(base) sib:~/kishan$ conda install tensorflow-gpu==1.12.0
WARNING: The conda.compat module is deprecated and will be removed in a future release.
Collecting package metadata: done
Solving environment: done

## Package Plan ##

  environment location: /home/sibridge/kishan/anaconda3

  added / updated specs:
    - tensorflow-gpu==1.12.0
```

Figure 2.14. TensorFlow Installation

- After installation, enter **Y** as shown in Figure 2.15.

```
wurlitzer      1.0.2-py37_0 --> 1.0.2-py36_0
xlrd           1.2.0-py37_0 --> 1.2.0-py36_0
xlwt           1.3.0-py37_0 --> 1.3.0-py36_0
zict           0.1.4-py37_0 --> 0.1.4-py36_0
zipp           0.3.3-py37_1 --> 0.3.3-py36_1

Proceed ([y]/n)? y_
```

Figure 2.15. TensorFlow Installation Confirmation

Figure 2.16 shows TensorFlow installation is completed.

```
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
(base) sib:~/kishan$ _
```

Figure 2.16. TensorFlow Installation Complete

2.2.5. Installing the Python Package

To install the Python package:

- Install Easydict by running the command below:

```
$ conda install -c conda-forge easydict
```

```
(base) sib:~/kishan$ conda install -c conda-forge easydict
Collecting package metadata: done
Solving environment: done

## Package Plan ##

  environment location: /home/sibridge/kishan/anaconda3

  added / updated specs:
    - easydict
```

Figure 2.17. Easydict Installation

- Install Easydict by running the command below:

```
$ conda install -c conda-forge librosa==0.6.0
```

```
(base)jay:~$ conda install -c conda-forge librosa==0.6.0
Collecting package metadata (current_repodata.json): done
Solving environment: failed with initial frozen solve. Retrying with flexible solve.
Collecting package metadata (repodata.json): done
Solving environment: done

## Package Plan ##

environment location: /home/jay/anaconda3

added / updated specs:
- librosa==0.6.0

certifi                                pkgs/main --> conda-forge
conda                                 pkgs/main --> conda-forge
openssl                               pkgs/main::openssl-1.1.1c-h7b6447c_1 --> conda-forge::openssl-1.1.1c-h516909a_0

Proceed ([y]/n)? y

Preparing transaction: done
Verifying transaction: done
Executing transaction: done
```

Figure 2.18. Anaconda Env – librosa Installation

Figure 2.18 shows the successful installation.

3. Preparing the Dataset

This chapter describes how to create a dataset using examples from Google Speech Command.

3.1. Selecting the Dataset

Below are two options of dataset for Key Phrase Detection models which can be used for training:

- Google Speech Command Dataset v0.01
- Google Speech Command Dataset v0.02

3.2. Downloading the Dataset

To download the dataset:

1. Go to the links below to download the Google Speech Command Dataset:
 - Google speech command dataset v0.01:
https://storage.cloud.google.com/download.tensorflow.org/data/speech_commands_v0.01.tar.gz
 - Google speech command dataset v0.02:
https://storage.cloud.google.com/download.tensorflow.org/data/speech_commands_v0.02.tar.gz
2. Convert the sample rate of 16 KHz .wav files to 8 KHz by running the commands below:

- a. Install the necessary dependency:

```
$ sudo apt-get install sox
```

- b. Convert the dataset:

```
$python convert-samplerate.py
```

```
earth:~/dataset$ python convert-samplerate.py --input_dir ./speech_commands_v0.02/ --output_dir ./speech_commands_v0.02_out --sample_rate 8000
sox WARN rate: rate clipped 1 samples; decrease volume?
sox WARN dither: dither clipped 1 samples; decrease volume?
sox WARN rate: rate clipped 71 samples; decrease volume?
sox WARN dither: dither clipped 61 samples; decrease volume?
sox WARN rate: rate clipped 13 samples; decrease volume?
sox WARN dither: dither clipped 11 samples; decrease volume?
sox WARN rate: rate clipped 61 samples; decrease volume?
sox WARN dither: dither clipped 57 samples; decrease volume?
```

Figure 3.1. Convert Sample Rate Script Execution

3.3. Data Augmentation

Dataset Augmentation is done by adding background noise in training-code itself. Background noises are used from the dataset's `_background_noise_` directory. Refer to README.md in the background noise directory for more information on how to add noise files.

4. Training the Machine

4.1. Training Code Structure

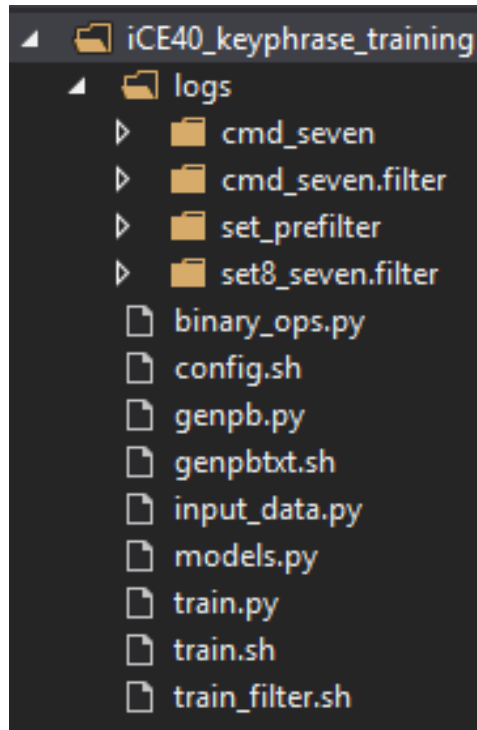


Figure 4.1. Training Code Directory Structure

4.2. Neural Network Architecture

4.2.1. Neural Network Architecture

This section provides information on the Convolution Neural Network Configuration of the Key Phrase Detection design.

Table 4.1. Key Phrase Detection Training Network Topology

Audio Input (8320)		
Freqconv	Conv1d-(64, 64)	<p>Conv1d basically performs convolution on 1d data and gives an output of 2D data. In this case,. the output is 64 x 64.</p> <p>Conv3 - # where:</p> <ul style="list-style-type: none"> Conv3 – 3 x 3 Convolution filter Kernel size # - The number of filters <p>For example, Conv3 - 8 = 8 3 x 3 convolution filter</p> <p>Batchnorm : Batch Normalization</p> <p>FC - # where:</p> <ul style="list-style-type: none"> FC – Fully connected layer # - The number of outputs
	Relu	
Reshape	Reshape- (64, 1, 64)	
Transpose	Transpose-(64, 64, 1)	
Fire 1	Conv3 – 16	
	Batchnorm	
	Relu	
	Maxpool	
Fire 2	Conv3 – 16	
	Batchnorm	
	Relu	
Fire 3	Conv3 – 20	
	Batchnorm	
	Relu	
	Maxpool	
Fire 4	Conv3 – 20	
	Batchnorm	
	Relu	
Fire 5	Conv3 – 20	
	Batchnorm	
	Relu	
	Maxpool	
Fire6	Conv3 – 24	
	Batchnorm	
	Relu	
	Maxpool	
Dropout	Dropout - 0.50	
fc4	FC – (2 + Num-keywords)	

- In [Table 4.1](#), Layer contains Convolution (conv), batch normalization (bn), Relu and dropout layers.
- Output of layer fc4 is number of classes in dataset along with silence and unknown considered as 2 keywords, so total number of outputs of fc4 layer is # of classes + 2.
- Layer information
 - Convolutional Layer

In general, the first layer in a CNN is always a convolutional layer. Each layer consists of number of filters (sometimes referred as kernels) which convolves with input layer/image and generates activation map (l.e. feature map). This filter is an array of numbers (the numbers are called weights or parameters). Each of these filters can be thought of as feature identifiers, like straight edges, simple colors, and curves & other high-level features. For example, the filters on the first layer convolve around the input image and *activate* (or compute high values) when the specific feature (for example, curve) it is looking for is in the input volume.

- **Relu (Activation layer)**

After each conv layer, it is convention to apply a nonlinear layer (or activation layer) immediately afterward. The purpose of this layer is to introduce nonlinearity to a system that basically has just been computing linear operations during the conv layers (just element wise multiplications and summations). In the past, nonlinear functions like tanh and sigmoid were used, but researchers found out that Relu layers work far better because the network is able to train a lot faster (because of the computational efficiency) without making a significant difference to the accuracy. The Relu layer applies the function $f(x) = \max(0, x)$ to all of the values in the input volume. In basic terms, this layer just changes all the negative activations to 0. This layer increases the nonlinear properties of the model and the overall network without affecting the receptive fields of the conv layer.
- **Pooling Layer**

After some Relu layers, programmers may choose to apply a pooling layer. It is also referred to as a down sampling layer. In this category, there are also several layer options, with Maxpooling being the most popular. This basically takes a filter (normally of size 2x2) and a stride of the same length. It then applies it to the input volume and outputs the maximum number in every sub region that the filter convolves around.

The intuitive reasoning behind this layer is that once we know that a specific feature is in the original input volume (there is a high activation value), its exact location is not as important as its relative location to the other features. As you can imagine, this layer drastically reduces the spatial dimension (the length and the width change but not the depth) of the input volume. This serves two main purposes. The first is that the number of parameters or weights is reduced by 75%, thus lessening the computation cost. Second is that it controls over fitting. This term refers to when a model is so tuned to the training examples that it is not able to generalize well for the validation and test sets. A symptom of over fitting is having a model that gets 100% or 99% on the training set, but only 50% on the test data.
- **Batchnorm Layer**

Batch normalization layer reduces the internal covariance shift. In order to train a neural network, we do some preprocessing to the input data. For example, we could normalize all data so that it resembles a normal distribution (that means, zero mean and a unitary variance). Reason being preventing the early saturation of non-linear activation functions like the sigmoid function, assuring that all input data is in the same range of values, etc.

But the problem appears in the intermediate layers because the distribution of the activations is constantly changing during training. This slows down the training process because each layer must learn to adapt themselves to a new distribution in every training step. This problem is known as internal covariate shift. Batch normalization layer forces the input of every layer to have approximately the same distribution in every training step by following below process during training time:

 - Calculate the mean and variance of the layers input.
 - Normalize the layer inputs using the previously calculated batch statistics.
 - Scales and shifts in order to obtain the output of the layer.

This makes the learning of layers in the network more independent of each other and allows you to be carefree about weight initialization, works as regularization in place of dropout, and other regularization techniques.
- **Drop-out layer**

Dropout layers have a very specific function in neural networks. After training, the weights of the network are so tuned to the training examples they are given that the network do not perform well when given new examples. The idea of dropout is simplistic in nature. This layer *drops out* a random set of activations in that layer by setting them to zero. It forces the network to be redundant. That means the network should be able to provide the right classification or output for a specific example even if some of the activations are dropped out. It makes sure that the network isn't getting too *fitted* to the training data and thus helps alleviate the over fitting problem. An important note is that this layer is only used during training, and not during test time.

- Fully connected layer**
This layer basically takes an input volume (whatever the output is of the conv or ReLU or pool layer preceding it) and outputs an N dimensional vector where N is the number of classes that the program must choose from.
- Quantization**
Quantization is a method to bring the neural network to a reasonable size, while also achieving high performance accuracy. This is especially important for on-device applications, where the memory size and number of computations are necessarily limited. Quantization for deep learning is the process of approximating a neural network that uses floating-point numbers by a neural network of low bit width numbers. This dramatically reduces both the memory requirement and computational cost of using neural networks.

The architecture above provides nonlinearities and preservation of dimension that help to improve the robustness of the network and control over fitting.

4.2.2. Key Phrase Detection Network Output

Key phrase detection network gives $N + 2$ values from last output node, where N is number of keywords trained. Two additional values represent *Silence* and *Unknown* as 2 keywords.

4.2.3. Training Code Overview

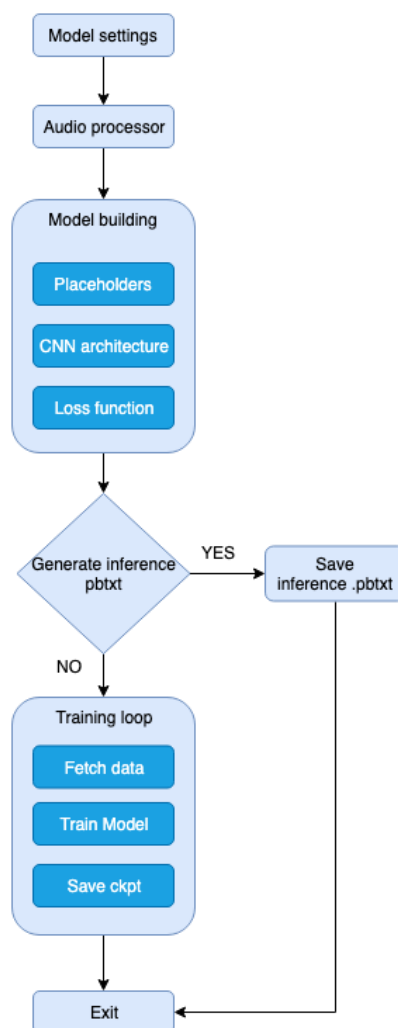


Figure 4.2. Training Code Flow Diagram

4.2.3.1. Model Settings

```
model_settings = models.prepare_model_settings(  
    len(input_data.prepare_words_list(FLAGS.wanted_words.split(','))),  
    FLAGS.sample_rate, FLAGS.clip_duration_ms, FLAGS.window_size_ms,  
    FLAGS.window_stride_ms, FLAGS.dct_coefficient_count)
```

Figure 4.3. Code Snippet – Model Settings

- *prepare_model_settings* function calculates and returns the dictionary containing all common settings for training from given parameters like *label_count* (# of keywords), *sample_rate*, *clip_duration_ms*, *window_size_ms*, *window_stride_ms* and *dct_coefficient_count*. More details about parameters are as follows:
 - *label_count* – Number of classes.
 - *sample_rate* – Number of audio samples per second.
 - *clip_duration_ms* – Length of each audio clip to be analyzed.
 - *window_size_ms* – Duration of frequency analysis window.
 - *window_stride_ms* – Length of move in time between frequency windows.
 - *dct_coefficient_count* – Number of frequency bins to use for analysis.

4.2.3.2. Audio Processor

```
audio_processor = input_data.AudioProcessor(  
    FLAGS.data_url, FLAGS.data_dir, FLAGS.silence_percentage,  
    FLAGS.unknown_percentage,  
    FLAGS.wanted_words.split(','), FLAGS.validation_percentage,  
    FLAGS.testing_percentage, model_settings)
```

Figure 4.4. Code Snippet – Audio Processor

```
class AudioProcessor(object):  
    """Handles loading, partitioning, and preparing audio training data."""  
  
    def __init__(self, data_url, data_dir, silence_percentage, unknown_percentage,  
                 wanted_words, validation_percentage, testing_percentage,  
                 model_settings):  
        self.data_dir = data_dir  
        self.maybe_download_and_extract_dataset(data_url, data_dir)  
        self.prepare_data_index(silence_percentage, unknown_percentage,  
                                wanted_words, validation_percentage,  
                                testing_percentage)  
        self.prepare_background_data()  
        self.prepare_processing_graph(model_settings)
```

Figure 4.5. Audio Processor Class Object

AudioProcessor function handles the loading, partitioning, and preparing of audio training data. It also downloads the dataset from given *data_url* in argument if dataset is not present.

```
desired_samples = model_settings['desired_samples']
self.wav_filename_placeholder_ = tf.placeholder(tf.string, [])
wav_loader = io_ops.read_file(self.wav_filename_placeholder_)
wav_decoder = contrib_audio.decode_wav(
    wav_loader, desired_channels=1, desired_samples=desired_samples)
# Allow the audio sample's volume to be adjusted.
self.foreground_volume_placeholder_ = tf.placeholder(tf.float32, [])
scaled_foreground = tf.multiply(wav_decoder.audio,
                                self.foreground_volume_placeholder_)
# Shift the sample's start position, and pad any gaps with zeros.
self.time_shift_padding_placeholder_ = tf.placeholder(tf.int32, [2, 2])
self.time_shift_offset_placeholder_ = tf.placeholder(tf.int32, [2])
padded_foreground = tf.pad(
    scaled_foreground,
    self.time_shift_padding_placeholder_,
    mode='CONSTANT')
sliced_foreground = tf.slice(padded_foreground,
                              self.time_shift_offset_placeholder_,
                              [desired_samples, -1])
# Mix in background noise.
self.background_data_placeholder_ = tf.placeholder(tf.float32,
                                                    [desired_samples, 1])
self.background_volume_placeholder_ = tf.placeholder(tf.float32, [])
background_mul = tf.multiply(self.background_data_placeholder_,
                              self.background_volume_placeholder_)
background_add = tf.add(background_mul, sliced_foreground)
background_clamp = tf.clip_by_value(background_add, -1.0, 1.0)
```

Figure 4.6. Code Snippet – Decoding the Audio, Scaling and Adding of Noise

4.2.3.3. Model Building

Placeholders

```
audio_input = tf.placeholder(
    tf.float32, [None, desired_samples], name='audio_input')
ground_truth_input = tf.placeholder(
    tf.int64, [None], name='groundtruth_input')
```

Figure 4.7. Code Snippet – Placeholders

- *audio_input* is placeholder to feed preprocessed audio input to network.
- *ground_truth_input* is placeholder to feed targeted labels of given batch. It is later used for evaluating training and optimization purpose.

CNN Architecture Generation

```
logits, dropout_prob, fingerprint_4d = models.create_model(  
    audio_input,  
    model_settings,  
    FLAGS.model_architecture,  
    is_training=True,  
    norm_binw=FLAGS.norm_binw,  
    downsample=FLAGS.downsample,  
    lock_prefilter=FLAGS.lock_prefilter,  
    add_prefilter_bias=FLAGS.prefilter_bias,  
    use_down_avgfilt=FLAGS.use_down_avgfilt)
```

Figure 4.8. Code Snippet – Create Model

create_model creates training graph or training model using given configuration. More details about flags are as follows:

- *audio_input* – TensorFlow node that gives audio feature map as output.
- *model_settings* – Dictionary of information about the model.
- *model_architecture* – The default architecture is *tinylvgg_conv* defined in *config.sh*.
- *is_training* – Whether the model is going to be used for training.

```
#####  
# VGG type  
def create_tinylvgg_conv_model(fingerprint_input, model_settings,  
                               is_training,  
                               filt_k=1,  
                               depthwise_conv1=False):  
    """Builds a convolutional model with low compute requirements.
```

Figure 4.9. Code Snippet – Create_tinylvgg_conv_model

Quantization and Network Model Configuration

```
if True: # 8b weight; 8b activation  
    fl_w_bin = 8  
    fl_a_bin = 8  
    ml_w_bin = 8  
    ml_a_bin = 8  
    ll_w_bin = 8  
    ll_a_bin = 16 # 16b results  
  
    min_rng = 0.0 # range of quantized activation  
    max_rng = 2.0
```

Figure 4.5. Code Snippet – Quantization Parameters Setting

```
fire1 = _fire_layer('fire1', fingerprint_4d, oc=depth[0], freeze=False, w_bin=fl_w_bin,
                    a_bin=fl_a_bin, min_rng=min_rng, max_rng=max_rng,
                    bias_on=bias_on, is_training=is_training)
fire2 = _fire_layer('fire2', fire1, oc=depth[1], freeze=False, w_bin=ml_w_bin, a_bin=ml_a_bin, pool_en=False,
                    min_rng=min_rng, max_rng=max_rng, bias_on=bias_on, is_training=is_training)
fire3 = _fire_layer('fire3', fire2, oc=depth[2], freeze=False, w_bin=ml_w_bin, a_bin=ml_a_bin,
                    min_rng=min_rng, max_rng=max_rng, bias_on=bias_on, is_training=is_training)
fire4 = _fire_layer('fire4', fire3, oc=depth[3], freeze=False, w_bin=ml_w_bin, a_bin=ml_a_bin, pool_en=False,
                    min_rng=min_rng, max_rng=max_rng, bias_on=bias_on, is_training=is_training)
fire5 = _fire_layer('fire5', fire4, oc=depth[4], freeze=False, w_bin=ml_w_bin, a_bin=ml_a_bin,
                    min_rng=min_rng, max_rng=max_rng, bias_on=bias_on, is_training=is_training)
fire6 = _fire_layer('fire6', fire5, oc=depth[5], freeze=False, w_bin=ml_w_bin, a_bin=ml_a_bin,
                    min_rng=min_rng, max_rng=max_rng, bias_on=bias_on, is_training=is_training)
fire_o = fire6
```

Figure 4.6. Code Snippet – Forward Graph Fire Layers

8-bit quantization is done on weights and activations in this model. Based on value of `w_bin` and `a_bin`, it is decided if quantization is to be done or not.

```
if w_bin == 8: # 8b quantization
    kernel_quant = lin_8b_quant(kernel)
    tf.summary.histogram('kernel_quant', kernel_quant)
    conv = tf.nn.conv2d(inputs, kernel_quant, [1, stride, stride, 1], padding=padding, name='convolution')

    if bias_on:
        biases = _variable_on_device('biases', [filters], bias_init, trainable=(not freeze))
        biases_quant = lin_8b_quant(biases)
        tf.summary.histogram('biases_quant', biases_quant)
        conv_bias = tf.nn.bias_add(conv, biases_quant, name='bias_add')
    else:
        conv_bias = conv
else: # 16b quantization
    conv = tf.nn.conv2d(inputs, kernel, [1, stride, stride, 1], padding=padding, name='convolution')
    if bias_on:
        biases = _variable_on_device('biases', [filters], bias_init, trainable=(not freeze))
        conv_bias = tf.nn.bias_add(conv, biases, name='bias_add')
    else:
        conv_bias = conv
```

Figure 4.7. Code Snippet – Convolution Quantization

4.2.3.4. Training Loop

Training Loop is divided in three segments:

- Fetching Data

```
train_audio, train_ground_truth = audio_processor.get_data(
    FLAGS.batch_size, 0, model_settings, FLAGS.background_frequency,
    FLAGS.background_volume, time_shift_samples, 'training', sess)
```

Figure 4.8. Code Snippet – Training Loop

Using `audio_processor` object, the code shown in Figure 4.8 fetches the next batch for training.

- Training Model

```
train_summary, train_accuracy, cross_entropy_value, _, _, fingerprint_4d_val = sess.run(
    [
        merged_summaries, evaluation_step, cross_entropy_mean, train_step,
        increment_global_step, fingerprint_4d
    ],
    feed_dict={
        audio_input: train_audio,
        ground_truth_input: train_ground_truth,
        learning_rate_input: learning_rate_value,
        dropout_prob: 0.5
    })
```

Figure 4.9. Code Snippet – Train Model

Using `session.run` code feeds next batch to network placeholders and evaluates desired operations mentioned in argument.

- Saving Checkpoints

```
if (training_step % FLAGS.save_step_interval == 0 or
    training_step == training_steps_max):
    checkpoint_path = os.path.join(FLAGS.train_dir,
                                   FLAGS.model_architecture + '.ckpt')
    tf.logging.info('Saving to "%s-%d"', checkpoint_path, training_step)
    saver.save(sess, checkpoint_path, global_step=training_step)
```

Figure 4.10. Code Snippet – Save Checkpoints

Trained Checkpoints are periodically saved based on flag specified `save_step_interval` at `checkpoint_path`.

4.2.4. Training from Scratch and/or Transfer Learning

Training of Key Phrase Detection model has two phases:

- Phase 1 – Training with all keywords available in dataset (Filter training).
- Phase 2 – Training for wanted keywords using the checkpoint of phase 1 model (Keyword training).

4.2.4.1. Modifying Training Keywords and Other Configurations

- To configure keywords or other parameters, modify `config.sh` under training directory. Default parameter values for Key Phrase Detection training are as shown in [Figure 4.11](#).


```

config.sh
1 #
2 # This file sets common params for key phrase training.
3 # Update log directory and dataset path if required.
4 #
5
6 # Update dataset path if needed
7 export DATA_DIR=./data/speech_commands_sAI2/
8
9 # Configure training log directory if needed
10 export FILTER_TRAIN_DIR=./logs/set8_seven.filter
11 export TRAIN_DIR=./logs/set_prefilter
12
13 export NETWORK=tinyvgg_conv
14 export TRAIN_OPT=--optimizer=Adam
15
16 # Network ID
17 export FILTER_NET_ID=cmd_seven.filter
18 export NET_ID=cmd_seven
19
20 # Keywords to train
21 export FILTER_TRAIN_KEYWORD="marvin,sheila,on,off,up,down,go,stop,left,right,yes,learn,visual,follow,\
22 no,cat,dog,bird,tree,house,bed,wow,happy,zero,one,two,three,four,five,six,seven,eight,nine,forward,backward"
23 export TRAIN_KEYWORD="seven,marvin,on,happy"
24
25 export COMMON_OPT="--model_architecture=$NETWORK \
26 --sample_rate=8000 \
27 --downsample=1 \
28 --no_prefilter_bias \
29 --clip_duration_ms=1040 \
30 --time_shift_ms=140.0 \
31 --window_size_ms=32.0 \
32 --window_stride_ms=16.0 \
33 --dct_coefficient_count=64 \
34 --background_volume=0.5"

```

Figure 4.11. Code Snippet – Config File config.sh

- Update `DATA_DIR` with the path to root directory of speech commands dataset.
- Update `FILTER_TRAIN_KEYWORD` with all available keywords in dataset.
- Update `TRAIN_KEYWORD` with all keywords that need to be trained by model.
- You can also configure additional parameters in `train_filter.sh` and `train.sh` if needed.

4.2.4.2. Filter training

After configuring (only if required) parameters mentioned in [Modifying Training Keywords and Other Configurations](#) section, run the script below to start filter training from scratch.

```
./train_filter.sh
```

```

earth:$ ./train_filter.sh
name: GeForce RTX 2080 Ti major: 7 minor: 5 memoryClockRate(GHz): 1.755
pciBusID: 0000:01:00.0
totalMemory: 10.73GiB freeMemory: 10.34GiB
2019-09-13 16:23:10.864151: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1512] Adding visible gpu devices: 0
[256, 1, 64]
[None, 8320, 1]
[None, 64, 64]
[None, 64, 64, 1]
INFO:tensorflow:Training from step: 1
2019-09-13 16:23:14.807664: I tensorflow/stream_executor/dso_loader.cc:152] successfully opened CUDA library libcublas.so.10.0 locally
INFO:tensorflow:Step #1: rate 0.010000, accuracy 3.0%, cross entropy 4.603645
INFO:tensorflow:Step #2: rate 0.010000, accuracy 7.0%, cross entropy 4.603001
INFO:tensorflow:Step #3: rate 0.010000, accuracy 5.0%, cross entropy 4.349138
INFO:tensorflow:Step #4: rate 0.010000, accuracy 7.0%, cross entropy 4.221001
INFO:tensorflow:Step #5: rate 0.010000, accuracy 5.0%, cross entropy 3.846268
INFO:tensorflow:Step #6: rate 0.010000, accuracy 7.0%, cross entropy 3.923666
INFO:tensorflow:Step #7: rate 0.010000, accuracy 5.0%, cross entropy 3.963864

```

Figure 4.12. Key Phrase Detection – Trigger Training with Default Options (Phase 1)

4.2.4.3. Keyword Training

- For phase 2 training, update the path of phase 1 checkpoint in *train.sh* by running the command below:
 - Make sure that below line refers to checkpoints generated by Phase 1 training.

```
TRAIN_OPT = "$TRAIN_OPT -set_prefilter=<path_to_traindir/tinyvgg_conv.ckpt-50000> --lock_prefilter"
```

- After configuring (only if required) parameters mentioned in [Modifying Training Keywords and Other Configurations](#) section, run the script below to start training from scratch.

```
./train.sh
```

```
earth:$ ./train.sh
name: GeForce RTX 2080 Ti major: 7 minor: 5 memoryClockRate(GHz): 1.755
pciBusID: 0000:01:00.0
totalMemory: 10.73GiB freeMemory: 10.34GiB
2019-09-13 16:25:19.800218: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1512] Adding visible gpu devices: 0
[256, 1, 64]
[None, 8320, 1]
[None, 64, 64]
[None, 64, 64, 1]
INFO:tensorflow:Restoring parameters from ./logs/set8_seven.filter/tinyvgg_conv.ckpt-50000
INFO:tensorflow:Training from step: 1
2019-09-13 16:25:23.611075: I tensorflow/stream_executor/dso_loader.cc:152] successfully opened CUDA library libcublas.so.10.0 locally
INFO:tensorflow:Step #1: rate 0.010000, accuracy 8.0%, cross entropy 3.645020
INFO:tensorflow:Step #2: rate 0.010000, accuracy 35.0%, cross entropy 2.330117
INFO:tensorflow:Step #3: rate 0.010000, accuracy 48.0%, cross entropy 2.245769
INFO:tensorflow:Step #4: rate 0.010000, accuracy 46.0%, cross entropy 2.318952
INFO:tensorflow:Step #5: rate 0.010000, accuracy 45.0%, cross entropy 2.491566
INFO:tensorflow:Step #6: rate 0.010000, accuracy 48.0%, cross entropy 2.219422
INFO:tensorflow:Step #7: rate 0.010000, accuracy 43.0%, cross entropy 2.098553
INFO:tensorflow:Step #8: rate 0.010000, accuracy 33.0%, cross entropy 2.271026
INFO:tensorflow:Step #9: rate 0.010000, accuracy 39.0%, cross entropy 1.720827
```

Figure 4.13. Key Phrase Detection – Training from Scratch (Phase 2)

4.2.4.4. Transfer Learning

- For transfer learning, *\$FILTER_TRAIN_DIR* and/or *\$TRAIN_DIR* should point to last iteration's log directory in *config.sh* and the latest checkpoint name should be updated in training script. New checkpoints are stored at the path given in *\$TRAIN_DIR* and/or *\$FILTER_TRAIN_DIR*.
- Modify below line in *train_filter.sh* and *train.sh* to specify checkpoints from where training should resume:

```
TRAIN_OPT="$TRAIN_OPT --start_checkpoint=$TRAIN_DIR/$NETWORK.ckpt-50000"
```

```
earth:$ ./train.sh
name: GeForce RTX 2080 Ti major: 7 minor: 5 memoryClockRate(GHz): 1.755
pciBusID: 0000:01:00.0
totalMemory: 10.73GiB freeMemory: 10.34GiB
2019-09-13 16:49:03.041945: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1512] Adding visible gpu devices: 0
[256, 1, 64]
[None, 8320, 1]
[None, 64, 64]
[None, 64, 64, 1]
INFO:tensorflow:Restoring parameters from ./logs/set_prefilter/tinyvgg_conv.ckpt-200
INFO:tensorflow:Training from step: 200
2019-09-13 16:49:07.087444: I tensorflow/stream_executor/dso_loader.cc:152] successfully opened CUDA library libcublas.so.10.0 locally
INFO:tensorflow:Step #200: rate 0.010000, accuracy 53.0%, cross entropy 1.226098
INFO:tensorflow:Saving to './logs/set_prefilter/tinyvgg_conv.ckpt-200'
INFO:tensorflow:Step #201: rate 0.001000, accuracy 67.0%, cross entropy 1.009409
INFO:tensorflow:Step #202: rate 0.001000, accuracy 69.0%, cross entropy 0.885104
INFO:tensorflow:Step #203: rate 0.001000, accuracy 76.0%, cross entropy 0.848068
INFO:tensorflow:Step #204: rate 0.001000, accuracy 67.0%, cross entropy 0.932998
INFO:tensorflow:Step #205: rate 0.001000, accuracy 70.0%, cross entropy 0.847289
INFO:tensorflow:Step #206: rate 0.001000, accuracy 58.0%, cross entropy 1.115180
```

Figure 4.14. Key Phrase Detection – Trigger Training with Transfer Learning

4.2.4.5. Training Status

- Training status can be checked in logs by observing different terminologies like cross entropy loss, learning rate, and training accuracy.

```
INFO:tensorflow:Step #792: rate 0.010000, accuracy 34.0%, cross entropy 2.272425
INFO:tensorflow:Step #793: rate 0.010000, accuracy 37.0%, cross entropy 2.290346
INFO:tensorflow:Step #794: rate 0.010000, accuracy 37.0%, cross entropy 2.079125
INFO:tensorflow:Step #795: rate 0.010000, accuracy 35.0%, cross entropy 2.284777
INFO:tensorflow:Step #796: rate 0.010000, accuracy 39.0%, cross entropy 2.117874
INFO:tensorflow:Step #797: rate 0.010000, accuracy 33.0%, cross entropy 2.304488
INFO:tensorflow:Step #798: rate 0.010000, accuracy 34.0%, cross entropy 2.265022
INFO:tensorflow:Step #799: rate 0.010000, accuracy 38.0%, cross entropy 2.370426
INFO:tensorflow:Step #800: rate 0.010000, accuracy 28.0%, cross entropy 2.225374
```

Figure 4.15. Key Phrase Detection – Training Logs

- You can use the TensorBoard utility for checking training status.

- Run the command below to start TensorBoard:

```
$ tensorboard --logdir=<log directory of training>
```

```
earth:~$ tensorboard --logdir logs/set8_seven/
TensorBoard 1.12.0 at http://earth:6006 (Press CTRL+C to quit)
```

Figure 4.16. TensorBoard – Launch

- This command provides the link `http://<name>:6006` which needs to be copied and open in any browser like Chrome, Firefox, and so on.

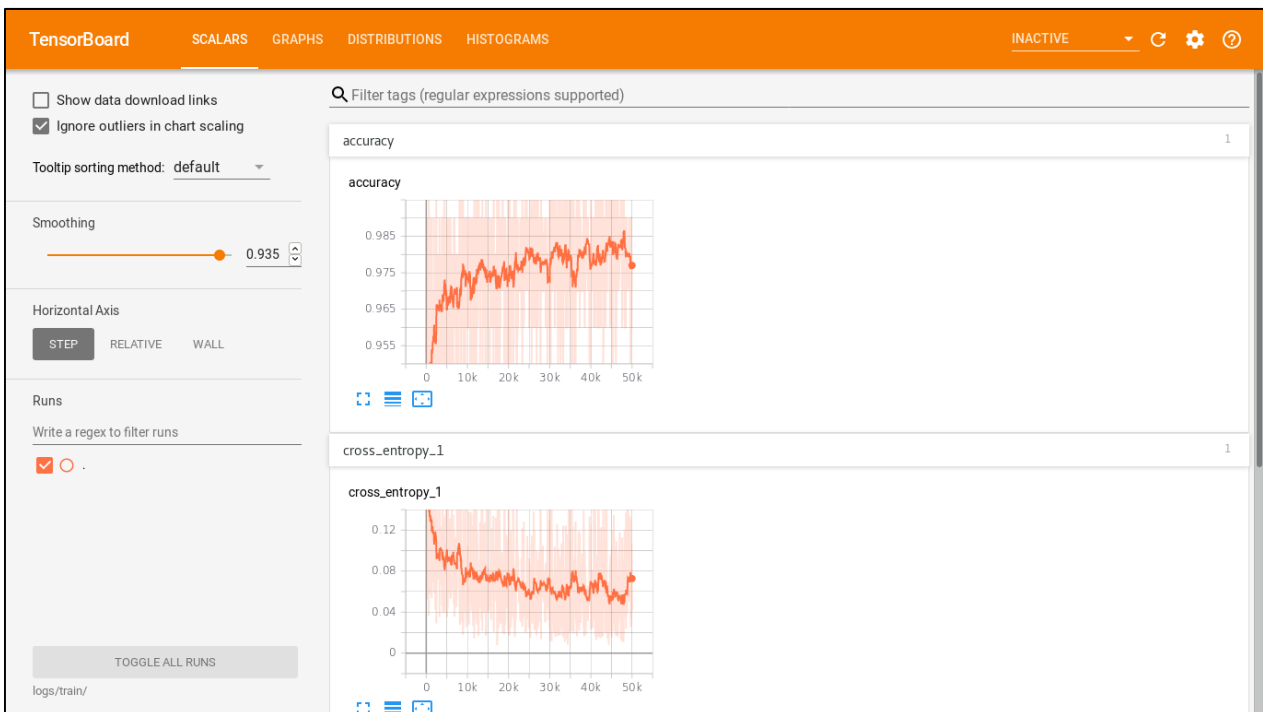


Figure 4.17. TensorBoard – Link Default Output in Browser

- Similarly, other graphs can be investigated from the available list.
- Check if the *checkpoint*, *data*, *meta* and *index* files are created at the log directory. These files are used for creating the frozen file (*.pb).




Figure 4.18. Checkpoint Storage Directory Structure

5. Creating Frozen File

This section describes the procedure for freezing the model, which is aligned with the Lattice SensAI tool. Perform the steps below to generate the frozen protobuf file:

```
earth:$ ./genpbtxt.sh
[256, 1, 64]
[None, 8320, 1]
[None, 64, 64]
[None, 64, 64, 1]
saved pbtxt for inference at log direcorry:./logs/set_prefilter
```

Figure 5.1. Generating .pbtxt For Inference




Figure 5.2. Generated .pbtxt for Inference

It generates the .pbtxt for inference under train log directory.

5.1. Generating the Frozen (.pb) File

Generate .pb file from latest checkpoint using below command from the training code's root directory.

```
$ python genpb.py --ckpt_dir <COMPLETE_PATH_TO_LOG_DIRECTORY>
```

```
earth:$ python genpb.py --ckpt_dir logs/set_prefilter/
using checkpoint :tinyvgg_conv.ckpt-500.meta
inputShape shape [None, None, None, None]
inputShape shapes [None, None, None, None]
output_shapes of input Node [None, None, None, None]
**TensorFlow**: can not locate input shape information at: audio_input
node to modify name: "audio_input"
```

Figure 5.3. Run genpb.py to Generate Inference .pb

genpb.py uses the generated .pbtxt and latest checkpoint in train directory to generate frozen .pb file.

Once the *genpb.py* is executed successfully, the log directory now contains the *<ckpt-prefix>_frozenforinference.pb* file as shown in Figure 5.4.

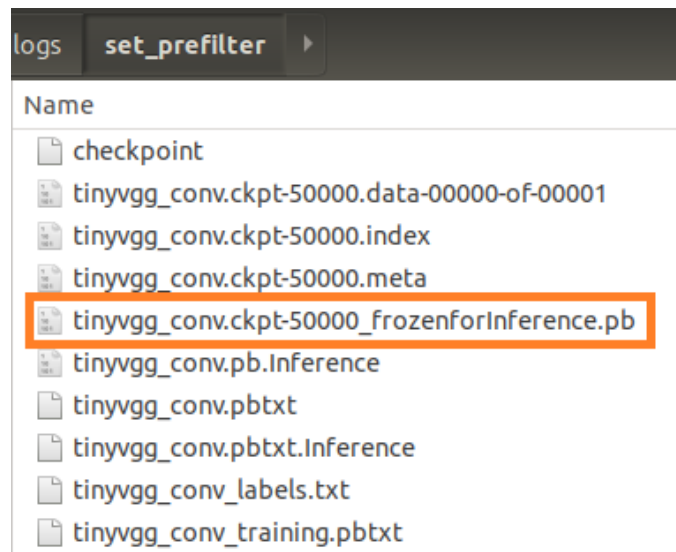


Figure 5.4. Frozen Inference .pb Output

6. Creating Binary File with SensAI

This chapter describes how to generate binary file using the Lattice SensAI version 2.1 program.



Figure 6.1. SensAI Home Screen

To create the project in SensAI tool:

1. Click **File > New**.
2. Enter the following settings:
 - **Project Name**
 - **Framework – TensorFlow**
 - **Class – CNN**
 - **Device – UltraPlus**
3. Click **Network File** and select the network (PB) file.




Figure 6.2. SensAI –Network File Selection

- Click **Image/Video/Audio Data** and select the image input file.


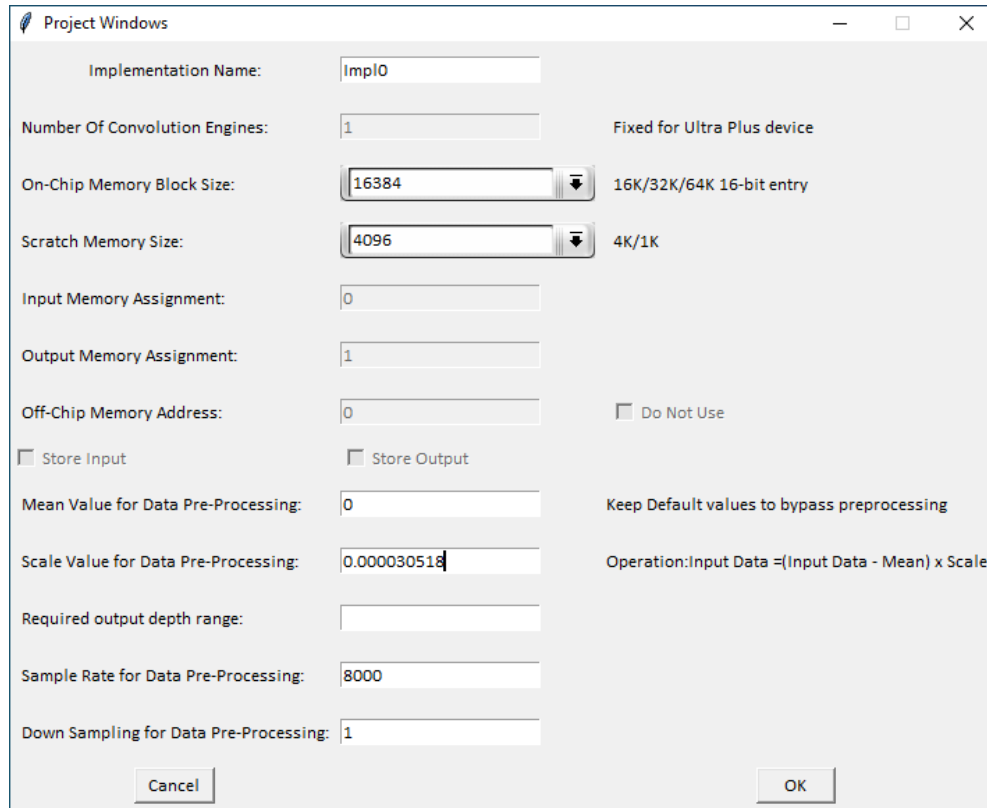


Figure 6.3. SensAI –Image Data File Selection

- Click **Next**.
- Configure your project settings as shown in [Figure 6.4](#).



Project Windows

Implementation Name:

Number Of Convolution Engines: Fixed for Ultra Plus device

On-Chip Memory Block Size: 16K/32K/64K 16-bit entry

Scratch Memory Size: 4K/1K

Input Memory Assignment:

Output Memory Assignment:

Off-Chip Memory Address: ☐ Do Not Use

☐ Store Input ☐ Store Output

Mean Value for Data Pre-Processing: Keep Default values to bypass preprocessing

Scale Value for Data Pre-Processing: Operation: $\text{Input Data} = (\text{Input Data} - \text{Mean}) \times \text{Scale}$

Required output depth range:

Sample Rate for Data Pre-Processing:

Down Sampling for Data Pre-Processing:

Figure 6.4. SensAI – Project Settings

7. Click **OK** to create project.
8. Double-click **Analyze**.

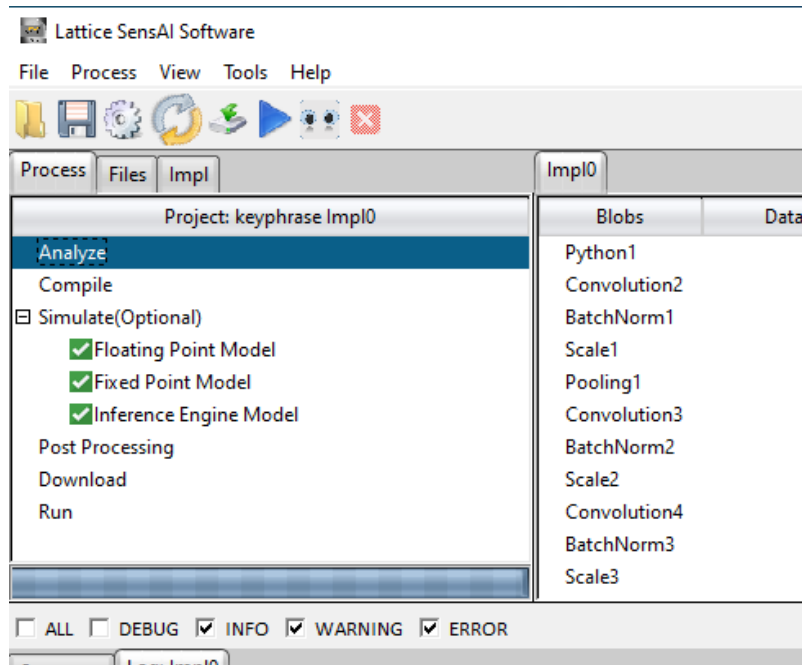


Figure 6.5. SensAI – Analyze Project

9. Double-click **Compile** to generate the Firmware file.

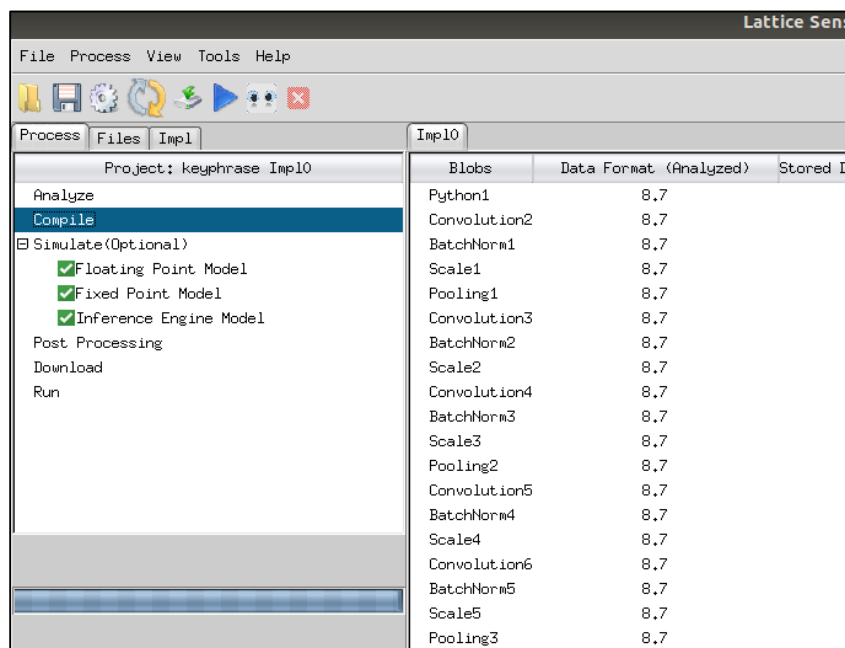


Figure 6.6. Compile Project

Firmware bin file location is displayed in the compilation log. Use the generated firmware and filter bin on the hardware for testing which gives you two output bin files:

- **<Project name>_filter.bin** – contains weights to generate spectrogram from audio input.
- **<Project name>.bin** – contains remaining neural network and weights.

7. Hardware Implementation

7.1. Top Level Information

7.1.1. Block Diagram

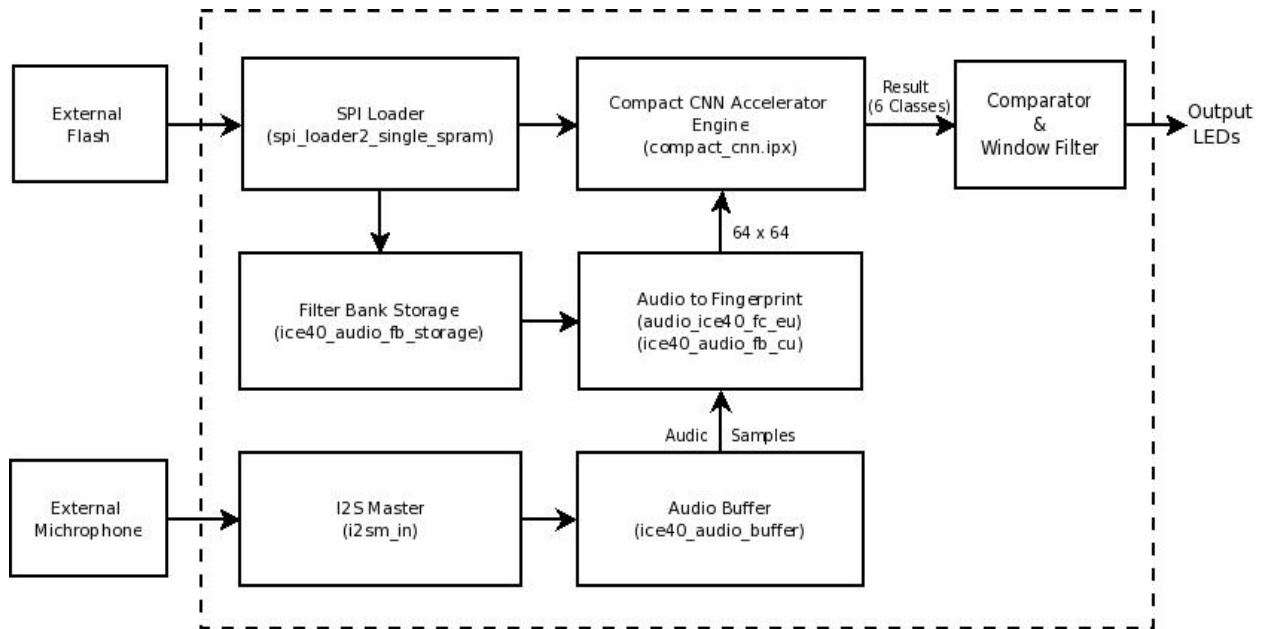


Figure 7.1. Top Level Block Diagram Key Phrase Detection iCE40

7.1.2. Overall Operational Flow

This section provides a brief idea about the data flow across ice40 Upduino2 Board.

- The I2S master module communicates with I2S microphone on Upduino2 board. It receives the serial audio data from the External Microphone and transports the data to the audio buffer module for storing purpose.
- The SPI loader module reads the external flash for two different files: the command sequence file and the filter bank storage file.
- The command sequence file is sent to the compact CNN IP core while the filter bank storage file is sent to the filter bank module. A filter bank storage module is used while converting audio data into a spectrogram-like picture for CNN input.
- The audio fingerprint module takes the stored audio data from the audio buffer along with the filter bank values from filter bank storage and creates an image of 64 x 64 which is then passed to CNN IP.
- As per command sequence received from SPI Loader, CNN IP performs operations on 64 x 64 input audio image and provides six values as output for six classes in following order:

Table 7.1. CNN Output Format

Index	0	1	2	3	4	5
Class	Silence	Unknown (No Keyword)	Seven (Keyword)	Marvin (Keyword)	On (Keyword)	Happy (Keyword)

- CNN outputs are passed to the audio post processing unit to compare all the six output values and generate difference value between 1st maximum value and 2nd maximum value (No Keyword). The index of 1st maximum value is verified to check whether it belongs to a valid keyword or not (valid Keyword Index values = 2, 3, 4, 5).
- The calculated difference value is compared with two Thresholds defined in Top module which is then used to drive the LEDs. The Detection LED logic uses another Threshold for Keyword Length which relates to the duration of the Keyword spoken. Currently the threshold values are as following:
 - Lower Threshold value = 1536 (Decimal)
 - Higher Threshold Value = 4096 (Decimal)
 - Keyword Length Threshold = 8 (Decimal)
- There are six LEDs with potential to turn on. According to the parameter EN_DEBUG_LED configuration, the demo has following two output representations.
 - Configuration 1 (EN_DEBUG_LED = 0)
 - LED D1 – It is ON if active audio is detected (including noise), and OFF when silence is detected.
 - LED D2 – OFF
 - LED D3 – OFF
 - LED D4 – OFF
 - LED D5 – OFF
 - LED D6 (Detection LED) – It is ON if any of the Keywords (Sheila, Marvin, On, Off) is detected and Keyword Length Threshold is crossed.
 - Configuration 2 (EN_DEBUG_LED = 1)

LED D1 and D6 are driven same as configuration 1 above. LED D2 to D5 is used for debugging information.

 - LED D1 – It is ON if active audio is detected (including noise), and OFF when silence is detected.
 - LED D2: – t has similar behavior as LED D6.
 - LED D3 – It is ON if any of the Keywords is detected with the calculated difference value greater than Higher Threshold.
 - LED D4 – It is ON if any of the Keywords is detected with the calculated difference value greater than Lower Threshold.
 - LED D5 – It is ON if any of the Keywords is detected with the calculated difference value lesser than both the Thresholds.
 - LED D6 (Detection LED) – It is ON if any of the Keywords (Sheila, Marvin, On, Off) is detected and Keyword Length Threshold is crossed.

7.1.3. Core Customization

Table 7.2. Core Parameters

Parameter	Default (Decimal)	Description
Configurable Parameters		
EN_DEBUG_LED	0	Enable additional LED detection for threshold value debugging. 1: D2 - D5 LEDs represents debug information for keyword detection with difference value in between certain threshold range. 0: D2 - D5 LEDs are always Off.

8. RTL Bitstream Generation

This section provides the procedure for creating your FPGA bitstream file using Lattice Radiant Software.

Note: This reference design includes a Compact CNN IP that requires a license to be able to generate a bitstream. Lattice provides a 30-day evaluation license for this IP for those who want to evaluate the IP and reference design. You can obtain an evaluation license from the Lattice website [Software Licensing](#) page.

Lattice Radiant software version 1.1 is required to generate a bitstream along with a software license patch. You can obtain the software patch file from the Lattice website through [Lattice Radiant 1.1 Software Patch](#).

To create the FPGA bitstream file:

1. Open the Lattice Radiant software.

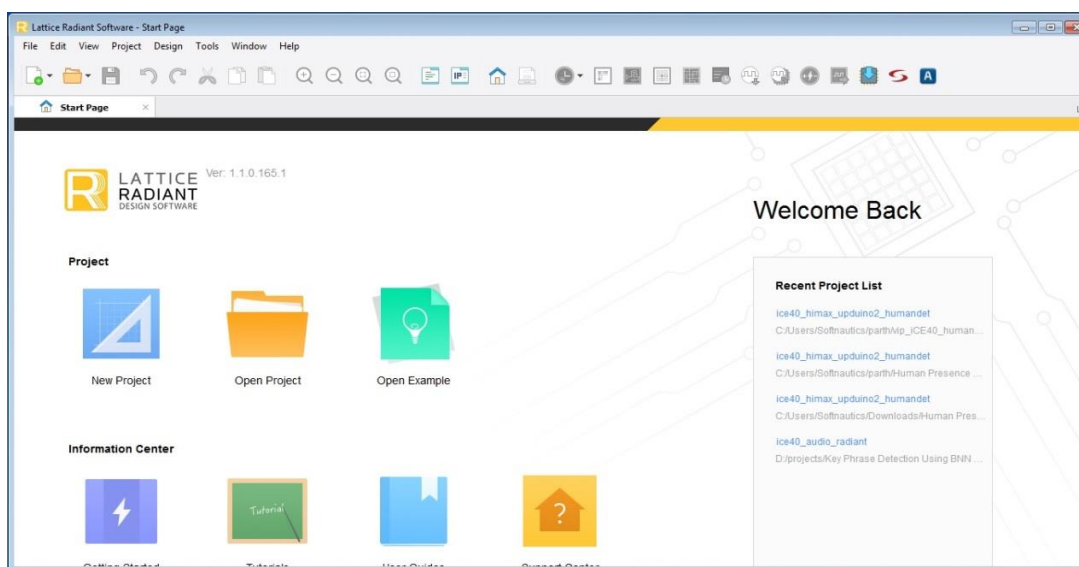


Figure 8.1. Lattice Radiant Software

2. Click **File > Open > Project**. Browse and open Radiant project file for iCE40 key phrase detection RTL.

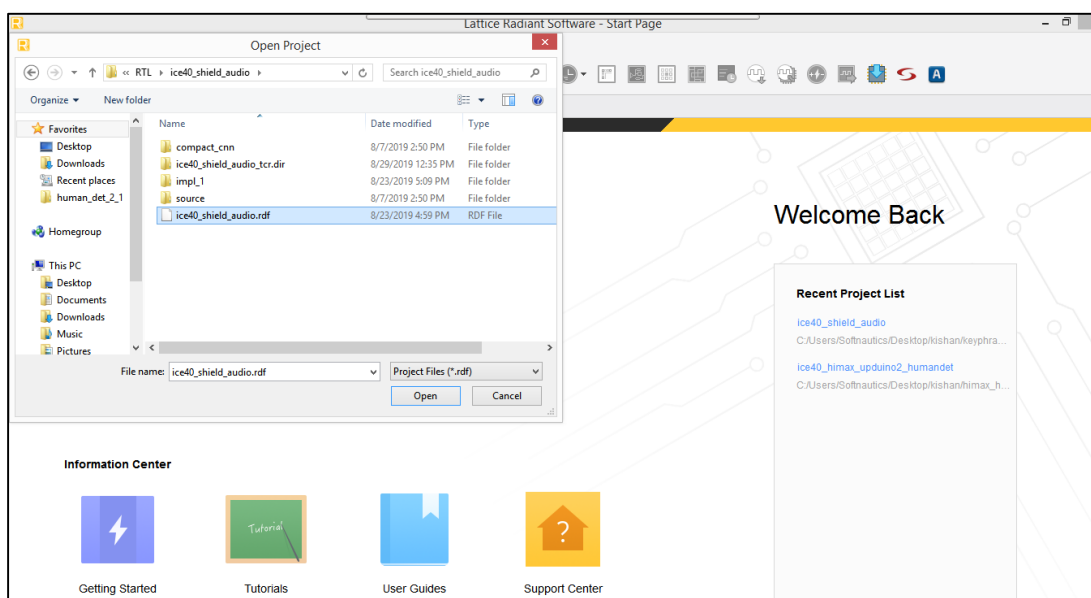


Figure 8.2. Lattice Radiant Software – Open Project

3. Click **Export Files** to generate the bit file.

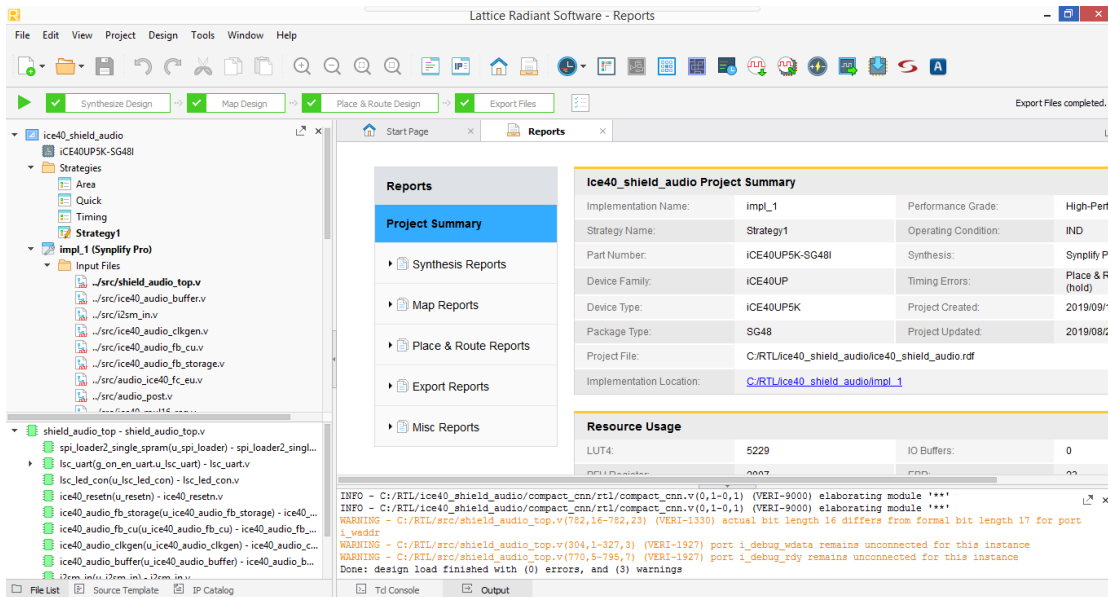


Figure 8.3. Lattice Radiant Software – Bitstream Generation

4. The **Export Reports** displays the generated bitstream as shown in Figure 8.4.

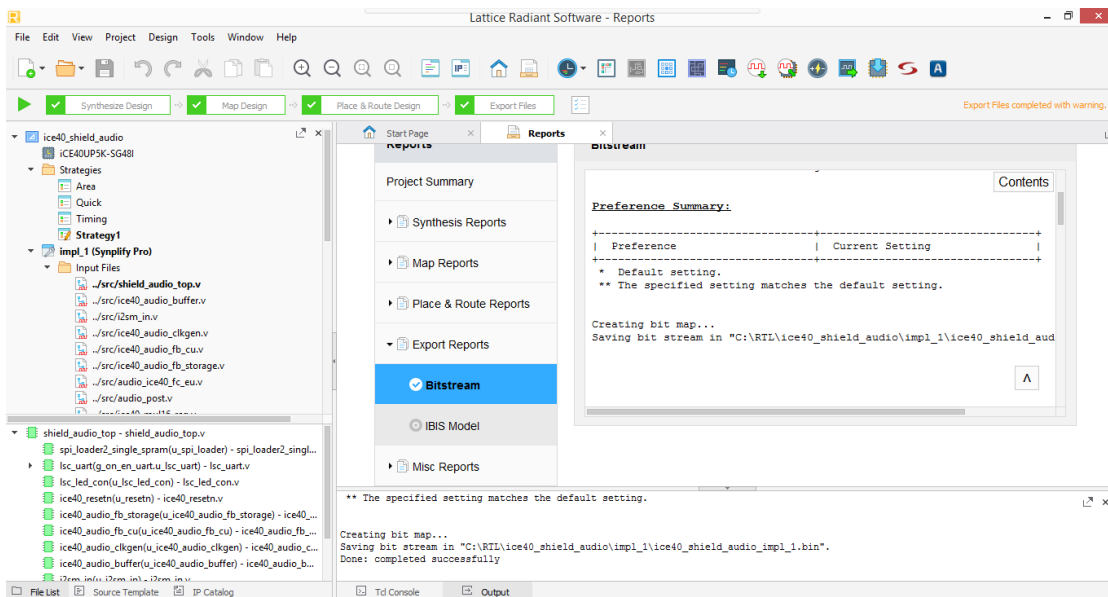


Figure 8.4. Lattice Radiant Software – Bitstream Generation Export Reports

9. Programming the Key Phrase Detection Demo

9.1. Functional Description

Figure 9.1 shows the diagram of the Key Phrase demo. The microphone captures audio and sends it to the iCE40 UltraPlus device. iCE40 UltraPlus uses the audio data with the firmware file from the external SPI Flash to determine the result.




Figure 9.1. iCE40 Key Phrase Demo Diagram

9.2. Programming Key Phrase Detection Demo on iCE40 SPI Flash

This section provides the procedure for programming the SPI Flash on the Himax HM01B0 UPduino Shield Board.

Two different files should be programmed into the SPI Flash. These files are programmed to the same SPI Flash, but at different addresses:

- Bitstream
- Firmware

To program the SPI Flash in Radiant Programmer:

1. Connect the Himax HM01B0 UPduino Shield board to the PC using a micro USB cable. Note that the USB connector on board is delicate, so handle it with care.
2. Start Radiant Programmer.
3. In the Radiant Programmer **Getting Started** dialog box, select **Create a new blank project**.
4. Click **OK**.

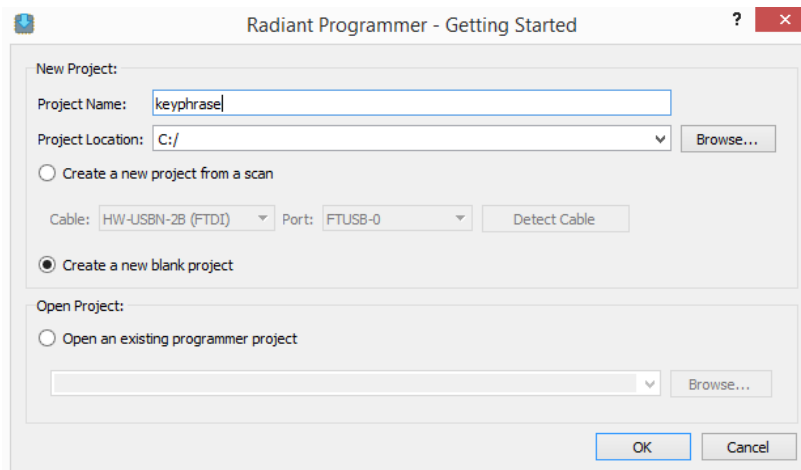


Figure 9.2. Radiant Programmer – Creating New Project

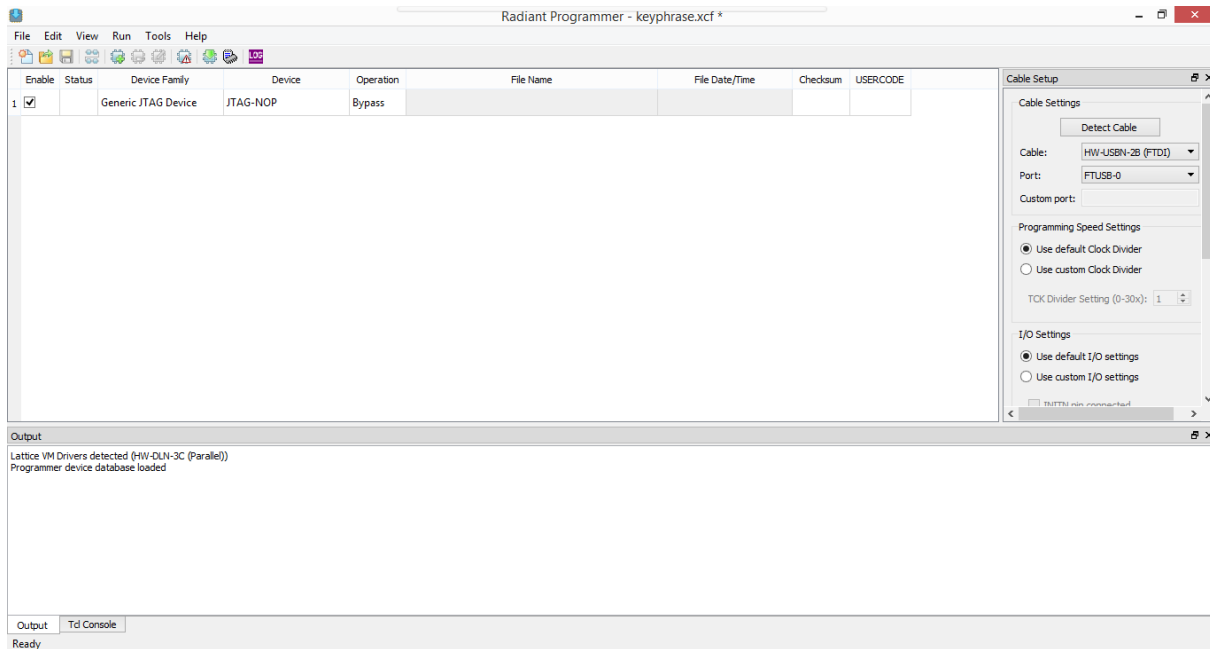



Figure 9.3. Radiant Programmer – Initial Project Window

- Initially, the .xcf have the option to add only one bin file. But since you need to program three bin files in case of key phrase demo, add two more devices by clicking the  button from the toolbar.
- Set Device Family to **iCE40 UltraPlus** and Device to **iCE40UP5K** for all three cases.

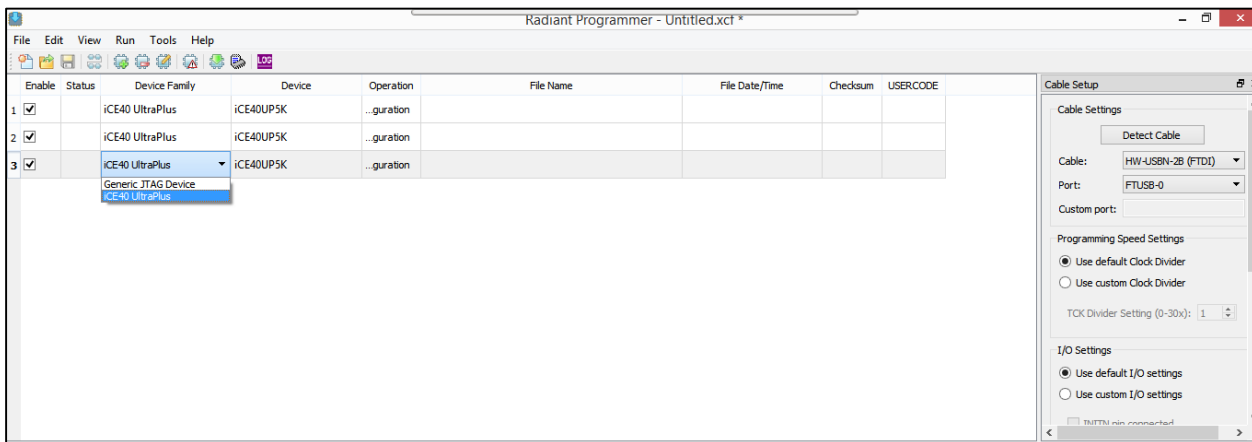


Figure 9.4. Radiant Programmer – iCE40 UltraPlus Device Family Selection

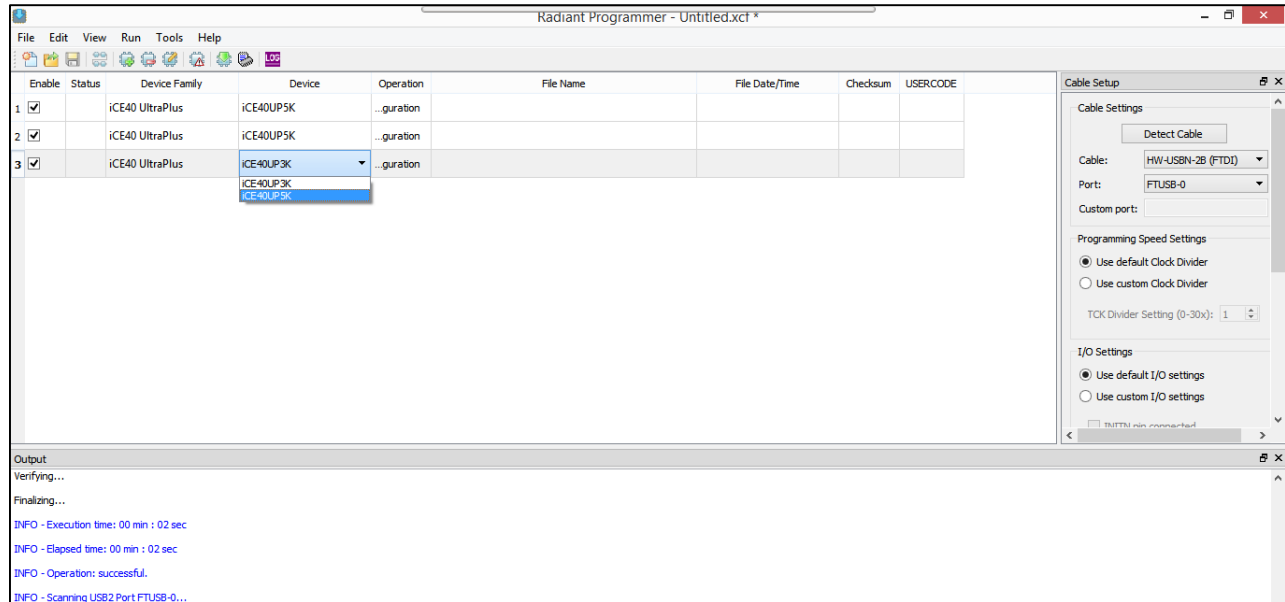


Figure 9.5. Radiant Programmer – iCE40 UltraPlus Device Selection

7. Click the iCE40 UltraPlus row and select **Edit > Device Properties**.
8. In the **Device Properties** dialog box, apply the settings below that are common to the two files to program.
Under Device Operation, select the options below:
 - **Target Memory – External SPI Flash Memory**
 - **Port Interface – SPI**
 - **Access Mode – Direct Programming**
 - **Operation – Erase, Program, Verify**
 Under SPI Flash Options, select the options below:
 - **Family – SPI Serial Flash**
 - **Vendor – Winbond**
 - **Device – W25Q32**
 - **Package – 8-pin SOIC**
9. To program the bitstream file, select the options below as shown in [Figure 9.6](#).
 - Under Programming Options, select the key phrase RTL bitstream file in Programming file.
 - Click Load from File to update the Data file size (Bytes) value.
 - Ensure that the following addresses are correct:
 - Start Address (Hex) – **0x00000000**
 - End Address (Hex) – **0x00010000**
10. Click **OK**.

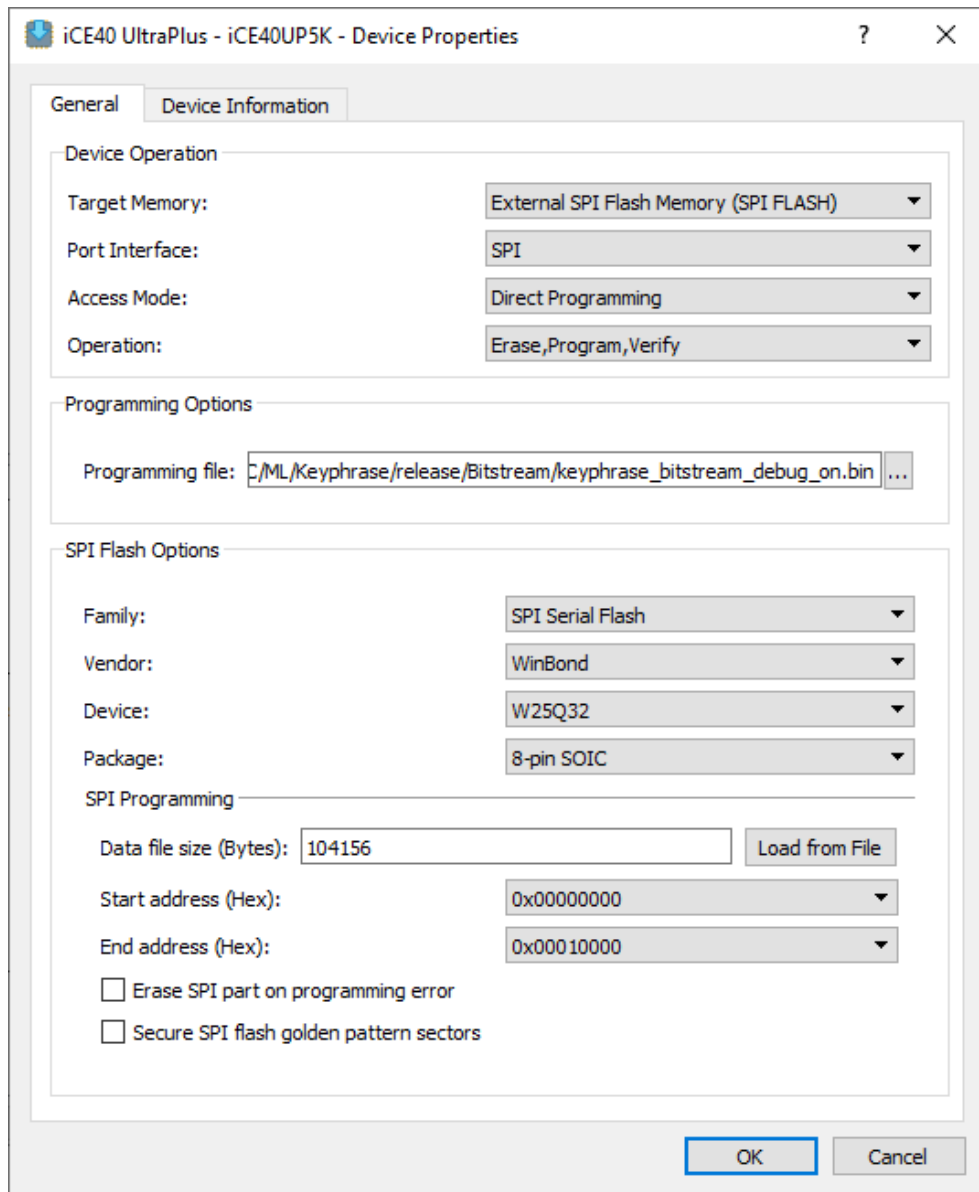


Figure 9.6. Radiant Programmer – Bitstream Flashing Settings

11. To program the filter binary firmware for generating the spectrogram, select the options below as shown in [Figure 9.7](#).
 - Under Programming Options, select the key phrase filter bin firmware generated by SensAI tool.
 - Click Load from File to update the Data file size (Bytes) value. Change **Data file size (Bytes)** value to **32768**.
 - Ensure that the following addresses are correct:
 - Start Address (Hex) – **0x00020000**
 - End Address (Hex) – **0x00030000**
12. Click OK.

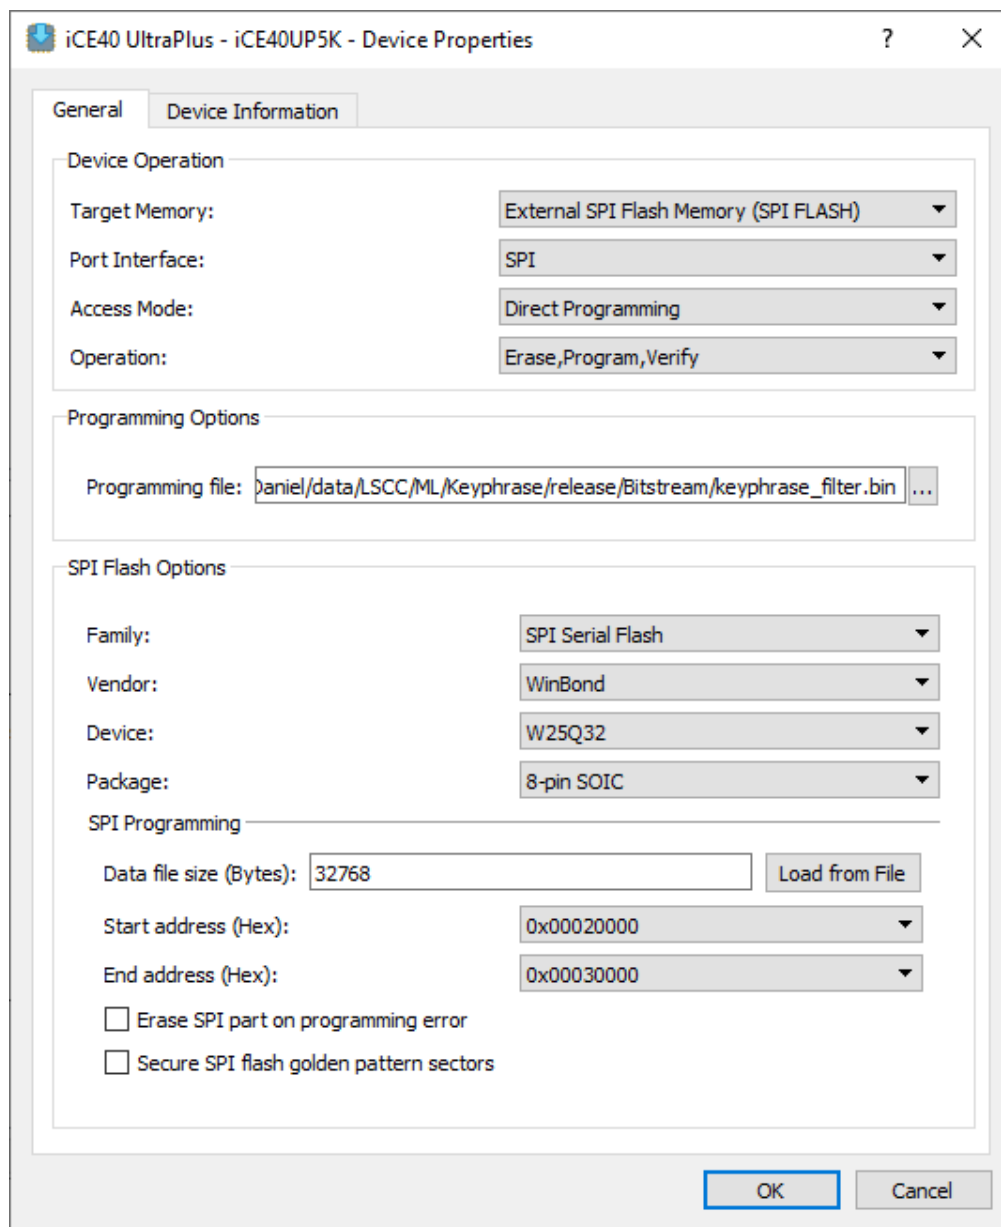


Figure 9.7. Radiant Programmer – Filter-Firmware Bin File Flashing Setting

13. To program firmware bin which contains model architecture, select the options below as shown in [Figure 9.8](#).
 - Under Programming Options, select the key phrase firmware binary file generated by SensAI tool.
 - Click Load from File to update the Data file size (Bytes) value. Change **Data file size (Bytes)** value to **31048** (In case model is changed, write the actual size of the file).
 - Ensure that the following addresses are correct:
 - Start Address (Hex) – **0x00030000**
 - End Address (Hex) – **0x00040000**
14. Click **OK**.

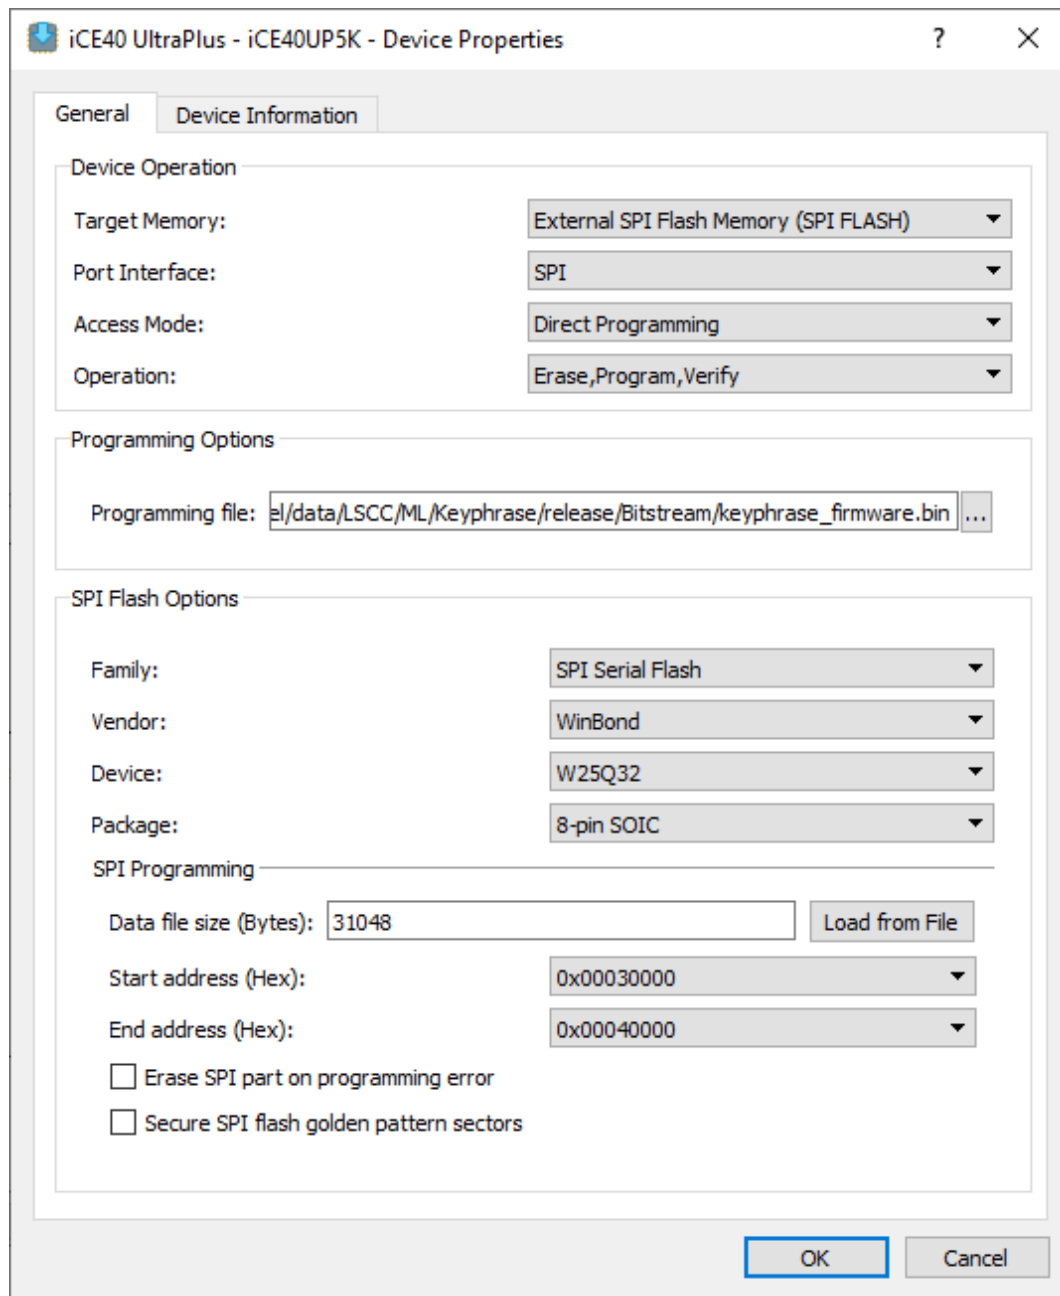


Figure 9.8. Radiant Programmer – Firmware Bin File Flashing Setting

15. In the main interface, click **Program Device** to program the binary file.

9.3. Run iCE40 Key Phrase Detection Demo on Hardware

To run the demo and observe results on the board:

1. Power ON the Himax HM01B0 UPduino Shield Board.
2. Speak the keyword in front of the board.
3. An LED light turns on if the keyword is detected. Refer to [Figure 9.9](#) for the LED information.




Figure 9.9. Camera and LED Location

- **D1** – It is ON if active audio is detected (including noise), and OFF when silence is detected.
- **D2-D4** – These are Debug LEDs. For more information, refer to [Overall Operational Flow](#) section.
- **D6** – This LED is turned on when the keyword is detected.

Technical Support Assistance

Submit a technical support case through www.latticesemi.com/techsupport.

Revision History

Revision 1.0, October 2019

Section	Change Summary
All	Initial release.



www.latticesemi.com